

Email-Driven Customer Insights — First Findings

Date: 4 June 2026 **Accounts analysed:** Artis (AM: Ehab Kamel) · The Property Agency (AM: Linda D'Arcy) **Approach:** Email interaction as the input signal; QuickBase sales data (won/lost, revenue) as the outcome. The question throughout: *which email behaviours go with winning, and which don't.*

A note on scope up front: this is two customers — a proof that the method works, not a validated rule across the book. Everything below should be read as "here is what the data shows for these two accounts," with the broader test still ahead.

The headline: there is no single "what wins" — it depends on the account

The most important finding is also the least expected. We assumed email behaviour would separate winning quotes from losing ones in a consistent way. It does — but only for one of the two accounts, and the *reason* is what matters.

Artis is a transactional account — high volume, frequent reorders, lower-ticket jobs. Here, email behaviour separates wins from losses clearly.

The Property Agency is a relationship account — fewer, larger deals, more contacts. Here, the same email signals are essentially flat. Winning isn't about how the relationship emails; it's about which contact carries it.

The implication for what we build: a single, global "winning-behaviour" model would blend these two opposite dynamics into noise and serve neither. The insight engine has to recognise the *type* of account before it can say what drives outcomes there.

Finding 1 — At transactional accounts, the *kind* of conversation predicts the win

For Artis, the sharpest signal isn't speed or volume — it's what the conversation is *about*.

- **Winning quotes** sit in production-oriented conversations: job approvals and artwork submissions. The customer is engaged in *making the thing*.
- **Losing quotes** sit in price-and-paperwork conversations: quote requests and payment queries. The customer is shopping or stuck on billing.

Supporting this, winning quotes also show faster sustained responses ($\approx 4.7\text{h}$ vs 9h average reply time) and more customer-initiated (rather than rep-initiated) threads — winning is associated with customer *pull*, not rep *push*.

One honest caveat: winning quotes also sit in busier email windows, but that may partly be because larger, live jobs naturally generate more email — so we treat conversation *type* (which is meaningful) as the primary signal and conversation *volume* as supporting, not the other way round.

For The Property Agency, none of these separate winning from losing quotes. That's not a weaker version of the Artis pattern — it's a genuinely different account dynamic, and it's why the account-type point above matters.

Finding 2 — Value concentrates in one contact, and it's not the busiest one

At both accounts, roughly half the won revenue flows through a single contact — and the highest-*volume* contact is usually not the highest-*value* one.

- **Artis:** `graphics@` converts $\sim 80\%$ of its quotes and drives the majority of won value. `accounts.pay@` appears on 2,000+ quotes but converts $\sim 4\%$ — it's a CC/admin address that gets cited constantly but barely engages. Being named on a quote is not the same as driving it.
- **The Property Agency:** the value runs through `ry@`, even though a different contact (`ck@`) sends more email. Volume contact $\neq$ value contact again.

The practical takeaway: relationship effort should be **value-weighted**, not **volume-weighted**. We can now identify the value contact per account — which is directly usable for prioritising who to engage.

What we checked but didn't include — and a finding about how books differ

We tested whether this analysis extends to other AMs, starting with Nic Doyle. His book behaves differently, and structurally so — which is itself worth knowing.

Three things separate Nic's accounts from Artis and The Property Agency:

- **Nearly half his won revenue falls outside the email-covered period** — older accounts that pre-date email sync, or accounts with no synced email at all. That's a coverage limit, not behaviour.

- **A large further share converts with no email near the quote** (~\$1.4M), even at accounts that email actively — typically reorders placed without a fresh thread.
- **His book is a long tail of ~1,700 accounts**, versus Artis and The Property Agency being single large trade relationships. A long tail of smaller, reorder-driven accounts structurally generates less email per quote.

The takeaway isn't that email is useless for Nic — it's that email-derived insight will have **materially lower coverage** for a book shaped like his, and any approach has to expect and model that rather than treat it as a fixable gap. It also tells us something useful in its own right: AMs' books differ enough in shape that the *right* insight method may differ by AM, not just by account — an extension of the account-type finding above.

Caveats (so the findings above can be read at face value)

- **Two customers.** Proof of method, not a validated rule. The next step generalises it.
 - **Email is a strong but proximity-based link.** We associate emails to quotes by contact and timing, which correlates well with outcomes but isn't a confirmed "this email is about this quote" link — fine for finding patterns, to be tightened for anything that needs exact attribution.
 - **Sentiment wasn't tested.** A numeric sentiment feature isn't built yet, so tone didn't enter this round; the conversation-type (intent) signal carried it instead.
 - **Mailbox coverage:** Artis is well-covered through Ehab's mailbox (recency slightly understated — his mailbox was ~2 weeks behind on sync). The Property Agency has ~21% of email outside Linda's named mailbox, a minor attribution caveat.
-

Where this goes next

The account-type finding is the bridge to scale. Rather than one model, the path is: classify accounts by type (transactional vs relationship), then test whether the Artis pattern holds across transactional accounts generally. That's what turns a two-customer proof into something usable across the book — and it's the foundation for the outreach use-case, where the goal becomes "engage *this* value contact, in *this* kind of conversation, at *this* account, because that's what wins here."